

Retrieval-augmented generation in three stages.

First, parse each document into text. Second, split the text into chunks small enough to embed and retrieve precisely. Third, embed each chunk into a vector and store it for similarity search.

At query time you embed the question, find the nearest chunks, and pass them to a language model as grounding context.

Warren models the first three stages as a pipeline of workers, one per stage, each scalable on its own.